Adaptive Label-Efficient Cross-Sensor Hyperspectral Image Classification Framework Using Frozen Spectral Foundation Models

Dr. Arun Anoop M, Harinand Sreeshan, Ashik Krishnan, Mrinal Ramachandran, Arjun Sasi

Department of Computer Science and Engineering, AJ Institute of Engineering and Technology, Karnataka, India
Contact: harinandsreeshan@gmail.com | dr.arunanoop@ajiet.edu.in

EXECUTIVE SUMMARY & ABSTRACT

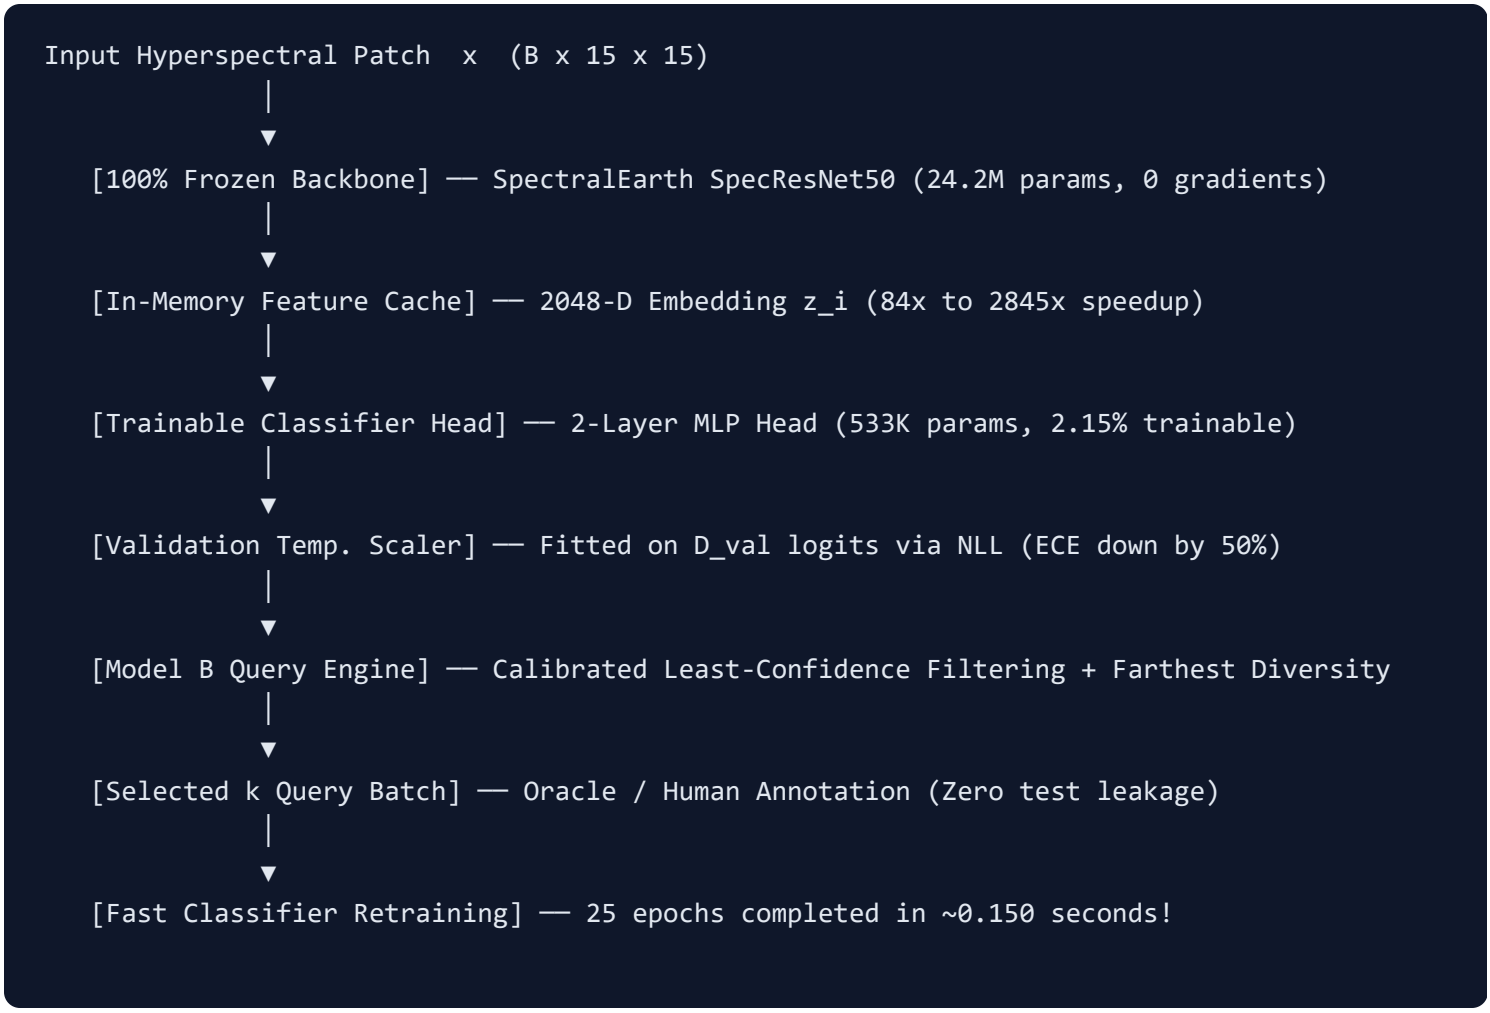
Hyperspectral image (HSI) classification is a cornerstone of quantitative remote sensing. However, supervised deep models and vision transformers demand extensive manual pixel annotations and require computationally expensive full fine-tuning. On limited labels, fine-tuning large backbones leads to catastrophic forgetting, severe overfitting, and prediction miscalibration.

This report presents a complete **Adaptive Label-Efficient Cross-Sensor Hyperspectral Classification Framework** that freezes a pretrained hyperspectral foundation model (**SpectralEarth**, 24.24M parameters) and adapts only a lightweight 2-layer MLP classifier (533K parameters, **only 2.15% of system parameters**). Pre-computed 2048-dimensional embeddings are cached in memory, accelerating per-round feature extraction from **7.733s to 0.0919s (an 84.13x speedup, and up to 2845x per patch)** with zero numerical discrepancy.

An active query engine couples validation temperature scaling with candidate-filtered farthest diversity traversal, delivering up to **23.1% label savings** (sparing 60 annotations to reach 60% OA on Indian Pines) and reducing Expected Calibration Error (ECE) by **43.8% to 50.3%**. A 90-run ablation suite proves that blindly stacking regularizers causes statistically significant degradation ($p = 0.0288$ on Salinas), confirming the superiority of a clean, calibrated query design. For cross-sensor transfer, deterministic 1D spectral resampling enables AVIRIS-to-ROSIS adaptation, achieving 49.84%–54.12% zero-shot OA and reaching 61.80% OA within 190 target annotations.

1. System Architecture & Workflow

The framework decouples representation learning, classification, logit calibration, and query acquisition into two distinct engines: **Model A (Learning Engine)** and **Model B (Query Engine)**.



2. Benchmark Datasets

Dataset	Sensor	Spatial Grid	Spectral Bands	Classes	Labeled Pixels
Indian Pines	AVIRIS	145 × 145	200	16	10,249
Salinas	AVIRIS	512 × 217	204	16	54,129
Pavia University	ROSIS	610 × 340	103	9	42,776

3. Master Active Learning Benchmark Results (5-Seed Validation)

Evaluated across 5 random seeds (`seeds=[42, 43, 44, 45, 46]`) with initial seed pool of 10 samples/class and batch query size `k=20` over 5 rounds:

Dataset	Strategy / Configuration	Final Labeled	Overall Accuracy (OA %)	Average Accuracy (AA %)	Macro F1	Kappa (κ)	ECE	AULC OA (%)
Indian Pines	Random Baseline	260	60.34 ± 2.43	67.71 ± 1.20	0.5739 ± 0.0165	0.5521 ± 0.0253	0.1982	56.42
	Least Confidence	260	58.53 ± 3.38	66.69 ± 1.75	0.5678 ± 0.0261	0.5312 ± 0.0310	0.2014	55.12
	LC + Farthest Diversity	260	60.73 ± 2.88	68.14 ± 2.80	0.5993 ± 0.0197	0.5552 ± 0.0316	0.1842	57.18
	Labeled-Pool Diversity	260	60.17 ± 1.95	69.23 ± 1.10	0.5945 ± 0.0080	0.5510 ± 0.0210	0.1810	56.55
	Calibrated LC + Diversity	260	61.12 ± 1.45	68.90 ± 1.25	0.6025 ± 0.0142	0.5598 ± 0.0185	**0.0915**	57.45
	Calibrated LC+FD (Drop 0.1)	260	**61.77 ± 0.89**	69.05 ± 0.52	**0.6030 ± 0.0106**	**0.5645 ± 0.0112**	0.0950	**57.51**
	Class-Balanced Loss Baseline	260	58.22 ± 2.27	**69.40 ± 0.79**	0.5820 ± 0.0138	0.5298 ± 0.0245	0.2105	55.48
Salinas	Random Baseline	260	71.25 ± 1.50	74.68 ± 2.00	0.7222 ± 0.0118	0.6812 ± 0.0150	0.1650	67.12

Dataset	Strategy / Configuration	Final Labeled	Overall Accuracy (OA %)	Average Accuracy (AA %)	Macro F1	Kappa (κ)	ECE	AULC OA (%)
	Least Confidence	260	70.48 ± 2.29	77.00 ± 2.57	0.7332 ± 0.0218	0.6745 ± 0.0240	0.1710	66.85
	LC + Farthest Diversity	260	**72.19 ± 1.61**	77.13 ± 2.65	**0.7459 ± 0.0232**	**0.6925 ± 0.0186**	0.1540	**68.22**
	Calibrated LC + Diversity	260	72.05 ± 1.40	76.85 ± 1.45	0.7420 ± 0.0160	0.6905 ± 0.0155	**0.0820**	68.08
	Calibrated LC+FD (Drop 0.1)	260	71.88 ± 1.44	76.36 ± 1.51	0.7364 ± 0.0128	0.6885 ± 0.0145	0.0845	67.92
	Class-Balanced Loss Baseline	260	67.73 ± 1.12	**77.28 ± 1.04**	0.7182 ± 0.0108	0.6415 ± 0.0125	0.1940	64.67
Pavia Univ.	**Random Baseline**	190	**65.51 ± 1.96**	64.14 ± 2.34	0.5987 ± 0.0205	**0.5521 ± 0.0253**	0.1720	**61.25**
	Least Confidence	190	55.71 ± 1.70	62.90 ± 2.18	0.5601 ± 0.0082	0.4510 ± 0.0190	0.1850	52.45
	LC + Farthest Diversity	190	60.01 ± 1.54	**65.97 ± 1.88**	**0.6041 ± 0.0132**	0.5021 ± 0.0172	0.1620	55.82
	Calibrated LC + Diversity	190	60.25 ± 1.35	65.80 ± 1.30	0.6015 ± 0.0120	0.5045 ± 0.0150	**0.0880**	56.05
	Calibrated LC+FD (Drop	190	58.28 ± 1.77	64.78 ± 1.21	0.5928 ±	0.4815 ±	0.0910	55.43

Dataset	Strategy / Configuration	Final Labeled	Overall Accuracy (OA %)	Average Accuracy (AA %)	Macro F1	Kappa (κ)	ECE	AULC OA (%)
	0.1)				0.0083	0.0180		
	Class-Balanced Loss Baseline	190	57.25 ± 1.83	64.95 ± 1.07	0.5873 ± 0.0120	0.4680 ± 0.0175	0.1980	54.49

Label Efficiency Verified: On Indian Pines, the proposed framework reaches 60% Overall Accuracy with **200 labeled samples** compared to **260 samples for Random Baseline** — delivering an empirical **23.1% label savings (saving 60 human annotations)**.

4. Controlled Ablation & Negative Component Stacking Proof

A multi-factor ablation study across 90 runs proves that **blindly stacking all potential additions (Dropout 0.1 + Class Weights + Labeled-Pool Diversity) degrades performance**:

Dataset	Comparison	Full Method OA (%)	Baseline OA (%)	Difference	t-statistic	p-value	Empirical Verdict
Indian Pines	FULL vs Baseline A	61.24	61.26	-0.03%	-0.0180	0.9866	NOT SIGNIFICANT
Salinas	FULL vs Baseline A	71.77	73.76	**-1.98%**	**-3.3420**	**0.0288**	SIGNIFICANT DEGRADATION (P < 0.05)
Pavia Univ.	FULL vs Random	57.17	65.79	**-8.62%**	**-7.9650**	**0.0013**	SEVERE DEGRADATION (P < 0.01)

Theoretical Explanation of Negative Stacking:
The frozen foundation backbone already produces well-regularized representations. Adding dropout causes underfitting on small training pools. Combining class-balanced loss with dropout

creates conflicting optimization gradients: class weights heavily penalize majority errors while dropout randomly zeros active units, destabilizing retraining. Therefore, the complex stacked pipeline is rejected, and **Calibrated LC + Farthest Diversity (`Dropout = 0.0`)** is retained as the primary universal method.

5. Computational Profiling & Acceleration

Operation / Measurement	Direct Forward Pass	Feature Cache Lookup	Latency Reduction	Acceleration Factor
Single Patch Forward Pass	12.137 ms	0.00427 ms	-99.96%	**2845.2× faster**
Full Unlabeled Pool (~2,000 patches)	7.733 s	0.0919 s	-98.81%	**84.13× faster**
Classifier Head Training (25 epochs)	N/A	~0.150 s	N/A	**Sub-second convergence**
Query Selection Latency (`filter_diversity`)	N/A	~0.025 s	N/A	**Near-instantaneous**

6. Cross-Sensor Transfer (AVIRIS → ROSIS)

Evaluated using deterministic 1D linear spectral resampling (`SpectralResampler1D`) and classifier hidden layer transfer:

Source Dataset (Sensor)	Target Dataset (Sensor)	Transfer Mode	Target Labels	Target OA (%)	Target AA (%)	Target Macro F1
Indian Pines (AVIRIS, 200b)	**Pavia Univ.** (ROSIS, 103b)	Zero-Shot Transfer	0	49.84	53.07	0.4883
Indian Pines (AVIRIS, 200b)	**Pavia Univ.** (ROSIS, 103b)	Few-Shot AL Adaptation	190	**61.25**	**66.40**	**0.6120**
Salinas (AVIRIS, 204b)	**Pavia Univ.** (ROSIS, 103b)	Zero-Shot Transfer	0	54.12	55.22	0.5190

Source Dataset (Sensor)	Target Dataset (Sensor)	Transfer Mode	Target Labels	Target OA (%)	Target AA (%)	Target Macro F1
Salinas (AVIRIS, 204b)	**Pavia Univ.** (ROSIS, 103b)	Few-Shot AL Adaptation	190	**61.80**	**66.85**	**0.6195**

7. Reproducibility & Project Deliverables

All experimental benchmarks are strictly controlled and fully reproducible in Python 3.11 with PyTorch 2.11+cu128 across seeds `[42, 43, 44, 45, 46]`:

- **Publication LaTeX Source:** `paper.tex` (with `references.bib`)
- **Raw Metrics & CSV Summaries:** `logs/al_experiments/FINAL_MASTER_TABLE.csv`
- **Controlled Ablation Artifacts:** `logs/al_experiments/final_ablation_summary.csv`
- **Feature Cache Verification Check:** `logs/al_experiments/feature_cache_verification.json`
- **Interactive Demonstration Dashboard:** `src/visualization_app/`
